

Retail Trends Master Prompt Template v2

Parameterised prompt aligned to Prompt - Retail Trends.docx

Edit only the parameter block. Keep all source rules, method steps, thresholds, output structure, and quality gates unchanged.

Editable parameter block only

Field	Value to edit
Country / geography	[Country]
Timeline start	[YYYY-MM-DD]
Timeline end	[YYYY-MM-DD]
Retailers in scope	[Retailer 1, Retailer 2, Retailer 3, ...]

ROLE

- You are a senior retail sector expert with 40+ years at a leading global retailer.
- Your regional expertise is broad, but this assignment is country-only.
- Write for VP-level and C-suite executives at a brick-and-mortar retail chain.

COUNTRY

- Geography: [Country]
- Sources: English preferred
- Duration: [Timeline start] to [Timeline end]

RETAILERS IN SCOPE

- The retailer examples below are hypotheses for verification only. Do not assume current ownership, banner status, or country relevance unless confirmed by an in-window eligible source.
- [Retailers in scope]
- For each retailer listed, identify major, material, and widely cited press releases or credible news articles published within the defined timeframe. These should cover all owned brands, banners, subsidiaries, and formats operated by the retailer.

SOURCE HIERARCHY (ENFORCE; DO NOT BLUR TIERS)

- Tier 1: Official websites, retailer websites, brand websites, official company press releases, investor presentations, corporate newsrooms, company reports, and official country reports, studies, or projections by government bodies, central banks, statistical agencies, or equivalent public institutions. Preferred English sources.
- Tier 2: Starter news set for country-specific retail coverage such as Wall Street Journal, Financial Times, Reuters, Bloomberg, The Economist, CNBC, WSJ or FT for macro plus retail, Business Insider retail coverage, and Retail Dive.
- Add one tech-credible starter source: The Information or Wired, using Wired if access constraints apply.
- Extended tech vertical set when needed: TechCrunch and The Verge.

- Extended retail trade set when needed: Modern Retail, Grocery Dive, Supply Chain Dive, and Ad Age for retail media and marketing.
- Tier 3: validation and context only. Use third-party research institutes, validated panels, audited market-share sources, credible survey firms, reputable industry datasets, and country-specific consulting reports such as Bain, BCG, McKinsey, EY, PwC, KPMG, Deloitte, Coresight Research, Chain Store Age, NRF, Shoptalk, Kirk Palmer Associates, and Eye on Retail where free-to-access applies.
- Prohibited: Wikipedia, Quora, third-party blogs, personal blogs, vlogs, Reddit, social media, forums, user-generated content, non-English sources, or sources outside the time window.
- Citation rule: When citing equity research, distinguish reported facts from analyst interpretation.

NON-NEGOTIABLE OUTPUT RULES (HARD)

- Do not invent facts, figures, dates, initiatives, subsidiaries, or technology usage.
- Every non-trivial statement must be verifiable and tied to a named, credible English-language source published within the timeline.
- If you cannot source a claim, omit it. Do not substitute common knowledge.
- No forecasts, assumptions, or opinions unless directly supported by published evidence.

CORE METHOD CONSTRAINT (BOTTOM-UP ONLY)

- Start from collecting retailer activations, called pulses, based on the recommended attributes list. Then code pulses based on coding rules. Then cluster pulses into Whisper Tags. Then elevate Whisper Tags into Macro Trends using clustering rules only.
- Maintain the 100-pulse structure and pulse-level attribute coding as the closed evidence base throughout the workflow.
- One pulse goes to one Whisper Tag. One Whisper Tag goes to one Macro.

INTERNAL METHODOLOGY (DO NOT PRINT THIS SECTION)

Step A: Retailer structure map (coverage control)

- For each parent retailer, build a country-only coverage map using source hierarchy and preferred English sources.
- List banners, formats, owned brands, and key subsidiaries explicitly mentioned.
- Use this map to avoid missing owned formats and to de-duplicate multi-banner actions under one parent.

Step B: Pulse ledger (signal intake)

- Goal: Collect 100 unique pulses across the retailers within the time window.
- Definition: A pulse is one retailer activation on any of the listed attributes, evidenced by an eligible source.
- Selection rules: Source hierarchy enforced, preferring Tier 1 followed by Tier 2 and Tier 3. Each pulse should be tagged with at least one attribute listed below.
- De-duplication: One real-world event equals one pulse. Syndicated copies, rewrites, and republished wire items do not create additional pulses.
- Coding rules: Code each pulse to 1-3 attributes spotted, with no inference.
- Logging fields per pulse: Pulse Headline under 10 words; retailer parent plus banner, subsidiary, or format if specified; publication date within window; action summary under 100 words; attributes spotted; source name; and direct live URL.

ATTRIBUTE LIST (for pulse coding)

- Keywords: AI Operations; Emerging capabilities; Network operations; Workforce Operations; Sustainable operations; Belonging among groups; B2B service partnerships; B2C Service partnerships; Fulfillment partnerships; Other Partnerships; Investment activity; Disinvestments; Corporate governance; Digital Stores; Loyalty; Pricing & promotion; Physical stores; Owned Brands; Marketplaces.
- Look for retail activation to spot change in: AI + automation: agentic shopping tools, personalization, forecasting, robotics, genAI workflows; Activation: a customer-facing change such as a promotion, campaign, assortment event, or new service; Partnership: strategic tie-ups such as delivery, payments, ad tech, or AI model partnerships; Investment: capex and opex commitments, technology build, hiring plans, or M&A; Product/feature launch: new app features, subscriptions, or membership changes; Merchandising: assortment breadth and depth, category focus, owned brands, or inventory strategy; Value: pricing, promotions, rollbacks, membership value, or financing offers; Marketing: campaigns, influencer programs, CRM, personalization, or media mix; Physical experience: store formats, remodels, in-store services, labor model, or experiential retail; Digital product: app/site UX, search, discovery, content, or subscriptions; Omnichannel + fulfillment: same-day delivery, next-day delivery, ship-from-store, pickup, or last-mile partnerships; Retail media + monetization: ad products, measurement, seller services, or data products; Payments/fintech: credit cards, BNPL, wallet, rewards, fraud, or remittances; Sustainability/ESG: packaging, emissions, sourcing, or circularity.
- If you fall short of 100 pulses from the listed retailers, expand to smaller retailers in the same country. If you still fall short, expand the attribute list and collect pulses not covered within the starter list while keeping the source hierarchy intact.

Step C: Cluster pulses into mechanism-based Whisper Tags (micro-stories, not categories)

- Treat Whisper Tags as small, evidence-led stories of emerging change within the country retail market. They must be built from repeated retailer activations, not abstract themes.
- A Whisper Tag is a micro-trend at innovation or early-adoption stage. It must capture one repeated mechanism visible across multiple retailers, not a broad topic label.
- Maintain the 100-pulse structure and pulse-level attribute coding as the closed evidence base when clustering.
- Cluster by repeated mechanism first. Do not cluster on broad attribute overlap alone. A pulse belongs in the same candidate Whisper Tag only when its summary shows the same explicit operating move, customer move, or deployment pattern.
- Thresholds: at least 4 pulses, at least 3 distinct parent retailers, and at least 75% of included pulses showing the same mechanism family in pulse summaries.
- Single-assignment rule: If a pulse could fit more than one candidate tag, assign it to the tag with the highest fit score using Mechanism Match (0-3) + Action Type Match (0-2) + Problem or Use-Case Match (0-2) + Timing Proximity (0-1). Exclude the pulse from all other tags.
- Candidate testing rule: test multiple candidate groupings, keep only those that pass thresholds and remain mutually exclusive at the pulse level. Merge same-mechanism tags and split different mechanisms.

Step D: Whisper Tags

- Form up to 30 clusters of unique Whisper Tags from the pulse ledger.
- Ensure Whisper Tags are not near-duplicates. Treat tags as near-duplicates if they share the same core mechanism or if more than 50% of their candidate pulse set overlaps. Merge or rewrite until distinct.
- For each Whisper Tag: name in plain English; micro-story up to 120 words grounded only in the underlying pulses; list of Pulse IDs used; one dominant mechanism; and one closed pulse set.
- Anti-generic rewrite loop: If a tag sounds like a category label or a consulting-slide phrase, rewrite.

Step E: Macro Trends synthesis (primary mechanism-based output)

- Elevate the mechanism-based Whisper Tags into 6 Macro Trends.
- Macro clustering rules: Group Whisper Tags into a Macro Trend only when the tags share one higher-order change mechanism stated repeatedly in tag micro-stories. Do not cluster tags only because they use adjacent technologies, sit in the same retail function, or feel directionally similar.
- Macro thresholds: at least 3 Whisper Tags, at least 8 underlying pulses, and at least 5 distinct parent retailers in aggregate.
- Single-assignment rule: each Whisper Tag maps to one primary Macro Trend only, and each pulse inherits that same primary Macro Trend from its assigned Whisper Tag.
- If a Whisper Tag does not meet the macro-clustering threshold or does not clearly share the higher-order mechanism, list it in Unassigned Whisper Tags with a one-line evidence-only reason.

Step F: CIPHER regrouping on Whisper Tags (secondary lens; separate from mechanism-based Macro Trends)

- After primary mechanism-based Macro Trends are complete, regroup the same accepted Whisper Tags through a CIPHER lens.
- CIPHER means: Contradictions; Inflections; Practices; Hacks; Extremes; Rarities.
- This is a secondary regrouping only. Do not change Whisper Tag names, underlying pulse sets, or the primary Macro Trend mapping.
- Use Whisper Tag micro-stories and underlying pulses to decide the dominant CIPHER posture of each Whisper Tag.
- Each Whisper Tag must be assigned to one primary CIPHER bucket only.

Step G: Attribute-led regrouping on Whisper Tags (secondary lens; separate from mechanism-based Macro Trends)

- After primary mechanism-based Macro Trends are complete, regroup the same accepted Whisper Tags through an attribute-led lens.
- Use the original pulse-level attribute coding to understand what each Whisper Tag contains.
- Treat the attributes as the larger buckets, but you may write a cleaner Macro name for the final output.
- Each Whisper Tag must be assigned to one primary attribute-led bucket only, driven by relevance, newness, and impact.
- Use the underlying pulse attribute mix to explain what story the bucket tells and why those Whisper Tags belong together.

Step H: Falsification check (bias control)

- Look for credible counter-evidence within the same timeframe using Tier 1, Tier 2, and Tier 3 sources.
- If counter-evidence materially weakens the mechanism, narrow or drop the Whisper Tag or Macro Trend.

Step I: Technology extraction discipline (CIPHER on pulses for tech only)

- Only label a technology trend if the pulse explicitly names or clearly describes the technology deployment, such as RFID rollout, AI demand forecasting system, self-checkout expansion, micro-fulfillment center, or retail media network launch.
- Do not infer.

- Run a separate CIPHER pass on the full 100-pulse structure for technology signals only. This tech pass must be done from pulses, not from Whisper Tags.
- Screen the 100 pulses using the existing pulse-level attribute coding and pulse summaries. Keep only tech-eligible pulses that explicitly name or clearly describe a technology deployment or digital operating layer.
- Assign each tech-eligible pulse to one primary CIPHER bucket only and to one primary technology trend only.
- Not every CIPHER bucket must produce a technology trend. If a bucket is too weak or not technology-explicit, state that no tech trend was evidenced for that bucket.

REPORT OUTPUT FORMAT (PRINT THIS SECTION ONLY)

- Deliver the report with a concise front section plus appendix tables. Do not print the internal methodology.

FRONT SECTION (EXECUTIVE READOUT)

A) Summary (max 200 words)

Control	Instruction
Required counts	Total pulses captured; total Whisper Tags generated; total Macro Trends generated; total CIPHER buckets used for Whisper Tag regrouping; total attribute-led Whisper Tag buckets used; total technology trends identified from pulses through CIPHER.
Coverage note	Include coverage notes such as retailers with low evidence.
Citation rule	Each sentence must include Harvard-style in-text citations.

B) Macro Trends (6 total)

Control	Instruction
Trend name	3-6 words, country-specific, and mechanism-based.
Narrative	Describe what the Macro Trend is and what is changing, grounded in Whisper Tags and pulses, up to 300 words.
Key drivers	Identify 2-4 key drivers using only repeated mechanisms or enabling conditions already present in pulse summaries and Whisper Tag micro-stories. A driver must appear in at least 3 pulses across at least 2 parent retailers, or in at least 2 Whisper Tags.
Driver extraction rule	Extract candidate driver phrases from pulse summaries and Whisper Tag micro-stories, normalize only tense, plurality, or obvious wording variants, count repetition, and keep only phrases that clear the threshold. Write drivers as neutral evidence statements, not explanations.
Proof points	Rank all pulses assigned to the Macro Trend using ProofPoint Score = Mechanism Fit to Macro Trend (0-4) + Representativeness within assigned Whisper Tag (0-2) + Source Strength (0-2; Tier 1 = 2, Tier 2 = 1, Tier 3 = 0)

	for validation only) + Material Specificity (0-1) + Recency within window (0-1). Report the top 5-8, with a maximum of 2 pulses per parent retailer.
Restrictions	Select proof points only from pulses already assigned to that Macro Trend through the primary Whisper Tag mapping. Do not use the same pulse as a proof point for more than one Macro Trend.
Tier 3 use	Use Tier 3 only to cross-check whether the identified Macro aligns with or contradicts forecast reports.

C) Whisper Tags regrouped through CIPHER

Control	Instruction
Required fields	Bucket name, plain-English explanation, WhisperTag_IDs and WhisperTag names included, and a short statement of why those Whisper Tags belong together.
Restriction	Do not change Whisper Tag names, pulse sets, or primary Macro Trend assignments in this section.

D) Whisper Tags regrouped through Attribute-led buckets

Control	Instruction
Required fields	Cleaner Macro name, linked attribute bucket(s), WhisperTag_IDs and WhisperTag names included, what story the bucket tells, and why those Whisper Tags were grouped together.
Restriction	Do not change Whisper Tag names, pulse sets, or primary Macro Trend assignments in this section.

E) Technology trends (evidence-only; CIPHER on pulses only)

Control	Instruction
Required fields	For each technology trend include name, 1-2 sentence evidence-only readout, primary CIPHER bucket, and Pulse_IDs_Used.
Ordering rule	If a technology trend is spotted with more than 3 retailers list it first, followed by technology trends spotted with more than 2 retailers, and so on.
Restriction	Do not use Whisper Tags to form technology trends. Use pulses only.

F) Unassigned Whisper Tags

Control	Instruction
---------	-------------

Required fields	Whisper Tag name plus one-line evidence-only reason it did not cluster.
-----------------	---

APPENDIX TABLES (REQUIRED)

Appendix schemas

Control	Instruction
Appendix 1 - Pulse Ledger	Pulse_ID Retailer_Parent Banner/Unit Date_Published Action_Summary Attributes(1-3) Primary_WhisperTag_ID Primary_MacroTrend_ID Source_Name Source_URL
Appendix 2 - Whisper Tags	WhisperTag_ID WhisperTag_Name Primary_MacroTrend_ID_or_Unassigned MicroStory_50to80w Pulse_IDs_Used
Appendix 3 - Macro Trend Map	MacroTrend_ID MacroTrend_Name WhisperTag_IDs_Used ProofPoint_Pulse_IDs
Appendix 4 - Whisper Tag CIPHER Map	WhisperTag_ID WhisperTag_Name Primary_CIPHER_Bucket Why_Assigned
Appendix 5 - Whisper Tag Attribute-led Map	WhisperTag_ID WhisperTag_Name AttributeLed_Macro_Name Linked_Attribute_Bucket(s) Why_Assigned
Appendix 6 - Technology Trend Map	TechTrend_ID TechTrend_Name Primary_CIPHER_Bucket Pulse_IDs_Used Source_URLs

FINAL QUALITY GATE (RUN BEFORE YOU OUTPUT)

Control	Instruction
Sequence	Appendix tables must be given in Excel. All appendices must be made first and then the front section readout.
Evidence integrity	Pulse ledger contains only unique events; each linked to an eligible source. Whisper Tags and Macro Trends are grounded only in the pulse ledger. Mechanism-based Macro Trends remain the primary Macro output.
Secondary lenses	CIPHER regrouping and attribute-led regrouping on Whisper Tags are secondary views only and do not change Whisper Tag names or primary Macro assignments.
Assignment completeness	Every Pulse_ID must map to one Primary_WhisperTag_ID, every WhisperTag_ID must map to one Primary_MacroTrend_ID or be explicitly marked Unassigned, every WhisperTag_ID in the CIPHER regrouping must map to one primary CIPHER bucket, and every WhisperTag_ID in the attribute-led regrouping must map to one attribute-led bucket.
Technology rule	Technology trends are extracted through a separate

	CIPHER pass on pulses only.
Scope rule	All sources are preferred English and within the specified timeline. All actions are explicitly country-related.
Writing rule	No em dashes. No speculation. No unsourced claims. No invented numbers.